

# Data Mining Assignment 2

Zhengyuan Zhu  
CSE5334 - Data Mining  
UNIVERSITY OF TEXAS AT ARLINGTON

April 4, 2020

## Problem 1: Non-parameteric density estimation

### Solution for 1.1

First, we set up all the parameters for generating the Gaussian random data where the  $\mu_1 = 5$  and  $\sigma_1 = 1$

```
1 data_5_1 = np.random.normal(5, 1, 1000)
```

Then, we write a function  $p = \text{mykde}(X, h, x)$  that performs kernel density estimation on data  $X$  with bandwidth  $h$  for 1-Dimensional data. The function take  $X$  as the samples that calculate the density of gaussian distribution on  $x$ . It traverses all the samples and add each sample's density one by one.

```
1 def mykde(X, h, x):  
2     density = np.zeros(len(X))  
3     for i in range(len(X)):  
4         density += np.exp(-(x-X[i])**2/(2*h**2))/(h*np.sqrt(2*np.pi))/len(  
5             X)
```

The figure below is the visualization of the data distribution:




Figure 1: Fake data with  $\mu_1, \sigma_1$

Then we do some visualizations of the probability density function. The figure below shows that the different densities with different bandwidth. The red line represents the ground truth distribution and the histogram represents the samples distribution. It shows that the distribution curve becomes more flat when bandwidths get larger and larger.




Figure 2: Histogram of X along with the different estimated probability density function

## Solution for 1.2

For the problem 1.2, the methodology essentially is the same as problem 1.1. The only different is that we add a new random normal distribution data with  $\mu_2 = 0$  and  $\sigma = 0, 2$ . We use the codes below for the visualization of results. And the conclusion is pretty much the same as problem 1.1. The larger bandwidth we got, the more flat probability density function in the figure.

```

1 fig = plt.figure(figsize=(12, 5))
2 ax = fig.add_subplot(111)
3 plt.hist(data_5_1, bins=30, edgecolor='black', density=True, label='
4         histogram of samples')
5
6 # Plot kernel function
7 x = np.arange(0, 10, 0.01)
8 y = normal(x, 5, 1)
9 d1 = mykde(data_5_1, 0.1, x)
10 d2 = mykde(data_5_1, 1, x)
11 d3 = mykde(data_5_1, 5, x)
12 d4 = mykde(data_5_1, 10, x)
13
14 # Plot the ground truth distribution and estimations
15 plt.plot(x, y, color='red', label='ground truth distribution')
16 ax.plot(x, d1, label='bandwidth=0.1', color='orange')
17 ax.plot(x, d2, label='bandwidth=1', color='green')
18 ax.plot(x, d3, label='bandwidth=5', color='black')
19 ax.plot(x, d4, label='bandwidth=10', color='purple')
20
21 plt.legend()
22 plt.xlabel('x-axis')
23 plt.ylabel('y-axis')

```




Figure 3: Histogram of X along with the different estimated probability density function

### Solution for 1.3

When we have 2 sets of 2-D Gaussian random data, we need to modify mykde function that can fit the 2-D gaussian distribution. I have tried a lot but failed, so that I decide to modify the sklearn vanilla Kernel Density and stack the data we generated to a grid form. The source codes of mykde2D is simple and easy to understand.

```

1 from sklearn.neighbors import KernelDensity
2
3 def mykde2D(x, y, bandwidth, xbins=100j, ybins=100j, **kwargs):
4     # create grid of sample locations (default: 100x100)
5     xx, yy = np.mgrid[x.min():x.max():xbins,
6                       y.min():y.max():ybins]
7
8     xy_sample = np.vstack([yy.ravel(), xx.ravel()]).T
9     xy_train = np.vstack([y, x]).T
10
11     kde_sk1 = KernelDensity(bandwidth=bandwidth, **kwargs)
12     kde_sk1.fit(xy_train)
13
14     # score_samples() returns the log-likelihood of the samples
15     z = np.exp(kde_sk1.score_samples(xy_sample))
16     return xx, yy, np.reshape(z, xx.shape)

```

The visualization of 2D gaussian distribution is a little bit tricky. The white dots represent the 2D gaussian distribution's samples. And I use a colored mesh to create a pseudocolor plot with a non-regular rectangular grid. The shade that cover the figure shows the  $\mu$  and  $\sigma$  as results of mykde2D. The results shows that when bandwidth is small(such as 0.1), the estimated distribution seems like to be "overfitting". And vice versa the estimated distribution become "underfitting" when bandwidth is large.




Figure 4: Test your function mykde on this data with  $h = 0.1$




Figure 5: Test your function mykde on this data with  $h = 1$




Figure 6: Test your function mykde on this data with  $h = 5$




Figure 7: Test your function mykde on this data with  $h = 10$

## Problem 2 Solution

### Solution 2.1

To solve the problem, the first thing is to generate 1000 training instances in two different classes (500 in each) from multi-variate normal distribution using the following parameters for each class. The source codes are below with samples.

```
1 mul1, sigma1 = [1, 0], [[1, 0.75], [0.75, 1]]
2 mul2, sigma2 = [0, 1], [[1, 0.75], [0.75, 1]]
3 size = 500
4
5 train_data2D_1 = np.random.multivariate_normal(mean=mul1, cov=sigma1, size
                                                =size)
6 train_data2D_1_label = np.zeros((size, 1))
7 print(train_data2D_1_label.shape)
8
```

```

9 train_data2D_2 = np.random.multivariate_normal(mean=mul2, cov=sigma2, size
                                                =size)
10 train_data2D_2_label = np.ones((size, 1))
11 X_train = np.vstack([train_data2D_1, train_data2D_2])
12 y_train = np.vstack([train_data2D_1_label, train_data2D_2_label])
13 print(X_train[:5], X_train[-5:])
14 print(y_train[:5], y_train[-5:])
15
16 # build test set
17 mul1, sigma1 = [1, 0], [[1, 0.75], [0.75, 1]]
18 mul2, sigma2 = [0, 1], [[1, 0.75], [0.75, 1]]
19 size = 500
20
21 test_data2D_1 = np.random.multivariate_normal(mean=mul1, cov=sigma1, size=
                                                size)
22 test_data2D_1_label = np.zeros((size, 1))
23 print(test_data2D_1_label.shape)
24
25 test_data2D_2 = np.random.multivariate_normal(mean=mul2, cov=sigma2, size=
                                                size)
26 test_data2D_2_label = np.ones((size, 1))
27 X_test = np.vstack([test_data2D_1, test_data2D_2])
28 y_test = np.vstack([test_data2D_1_label, test_data2D_2_label])
29 print(X_test[:5], X_test[-5:])
30 print(y_test[:5], y_test[-5:])
31 -----
32 (500, 1)
33 [[-0.19002878 -0.14795651]
34  [-0.27783423 -1.46171299]
35  [ 0.78327396  0.39920896]
36  [ 0.57195261 -0.04425206]
37  [ 1.0097563  -1.48342926]] [[ 1.0917887   2.14877639]
38  [ 0.20435192  1.65422036]
39  [-1.49104023  0.3028843 ]
40  [-0.48226042  0.62071859]
41  [-0.26311759 -0.31374102]]
42 [0.]
43 [0.]
44 [0.]
45 [0.]
46 [0.]] [[1.]
47  [1.]
48  [1.]
49  [1.]
50  [1.]]
51
52 (500, 1)
53 [[ 1.54609753 -0.01456578]
54  [ 1.49064799  0.83087897]
55  [-0.59881474 -1.20925996]
56  [ 0.34422204 -0.58984347]
57  [ 0.42700309 -1.39169054]] [[-0.21408786  0.96185428]
58  [ 0.56305549  1.13239575]
59  [ 2.30235643  3.27293605]

```

```

60 [ 1.5609519  1.07252668]
61 [-1.14839655  0.92135342]]
62 [0.]
63 [0.]
64 [0.]
65 [0.]
66 [0.]] [[1.]
67 [1.]
68 [1.]
69 [1.]
70 [1.]]

```

Then I implement your Naive Bayes Classifier `[pred, posterior, err] = myNB(X,Y,X_test,Y_test)` whose inputs are the training data **X**, labels **Y** for **X**, testing data **X\_test** and labels **Y\_test** for **X\_test** and returns predicted labels **pred**, posterior probability **posterior** with which the prediction was made and error rate **err**.

```

1 import pandas as pd
2 import numpy as np
3 print(pd.__version__)
4
5 def myNB(X, y, X_test, y_test):
6     def norm(x, mean, std):
7         return 1/(np.sqrt(2*np.pi)*std)*np.exp(-(x-mean)**2/2*(std**2))
8
9     X_train, y_train = pd.DataFrame(X, columns=['Feature1', 'Feature2']),
10                        pd.DataFrame(y, columns=['Class'])
11
12     X_test, y_test = pd.DataFrame(X_test, columns=['Feature1', 'Feature2']),
13                     pd.DataFrame(y_test, columns=['Class'])
14
15     y_unique = y_train['Class'].unique()
16     train_set = pd.concat([X_train, y_train], axis=1)
17     prior = np.zeros(len(y_unique))
18     conditional = np.reshape(np.zeros(len(y_unique)*len(X_train.columns)*2),
19                             (len(y_unique), len(X_train.columns), 2))
20
21     for i in range(0,len(y_unique)):
22         prior[i]=(sum(y_train['Class']==y_unique[i])+1)/(len(y_train['Class'])+len(y_unique))
23
24     # print("The prior probability of each class is: ", prior)
25
26     for i, h in enumerate(X_train.columns.values.tolist()):
27         for j in range(0,len(y_unique)):
28             class_feature = train_set[h].loc[(train_set['Class']==y_unique[j])]
29
30             mean = np.mean(class_feature)
31             var = np.std(class_feature)
32             conditional[i][j] = [mean, var]
33
34     # print("mean and standard variance of current feature is: ",
35           # h, j, mean, var)
36
37     # print(conditional)

```

```

29 #     print("Prior distribution is: ", prior)
30 pred_probs = []
31 for idx, row in X_test.iterrows():
32     probs = []
33     for cIdx, pri in enumerate(prior): # class 0 1
34         for fIdx, feat in enumerate(row): # feature 0 1
35             pri *= norm(feat, conditional[fIdx][cIdx][0], conditional[
36                                     fIdx][cIdx][1])
37         probs.append(pri)
38     pred_probs.append(probs)
39     pred = [np.argmax(p) for p in pred_probs]
40 # calculate error rate
41 #     print(np.array(pred).shape)
42 #     print(y_test.to_numpy().shape)
43 err = np.mean(np.array(pred) != y_test.to_numpy().squeeze())
44 #     print(err)
45
46     return pred, conditional, err
47
48 pred, posterior, err = myNB(X_train, y_train, X_test, y_test)

```

For the experiments part, I perform the experiments 10 times and take an average of error rate. The source codes and results are shown below.

```

1 # perform the experiments 10 times
2 avg_err, run_time = 0, 10
3 for i in range(run_time):
4     train_data2D_1 = np.random.multivariate_normal(mean=mul1, cov=sigma1,
5                                                     size=size)
6     train_data2D_1_label = np.zeros((size, 1))
7
8     train_data2D_2 = np.random.multivariate_normal(mean=mul2, cov=sigma2,
9                                                     size=size)
10    train_data2D_2_label = np.ones((size, 1))
11    X_train = np.vstack([train_data2D_1, train_data2D_2])
12    y_train = np.vstack([train_data2D_1_label, train_data2D_2_label])
13
14    test_data2D_1 = np.random.multivariate_normal(mean=mul1, cov=sigma1,
15                                                  size=size)
16    test_data2D_1_label = np.zeros((size, 1))
17
18    test_data2D_2 = np.random.multivariate_normal(mean=mul2, cov=sigma2,
19                                                  size=size)
20    test_data2D_2_label = np.ones((size, 1))
21    X_test = np.vstack([test_data2D_1, test_data2D_2])
22    y_test = np.vstack([test_data2D_1_label, test_data2D_2_label])
23
24    pred, posterior, err = myNB(X_train, y_train, X_test, y_test)
25    print("current time error: ", err)
26    avg_err += err/run_time
27 print("The average error: ", avg_err)
28 -----
29 current time error: 0.117
30 current time error: 0.057

```

```

27 current time error: 0.101
28 current time error: 0.061
29 current time error: 0.086
30 current time error: 0.08
31 current time error: 0.079
32 current time error: 0.08
33 current time error: 0.071
34 current time error: 0.09
35 The average error: 0.08219999999999998

```

I have performed prediction on the testing data with source code. The accuracy, precision, recall, confusion matrix as well as a scatter plot of data points whose labels are color coded are shown below.

```

1 from sklearn import metrics
2 import matplotlib.pyplot as plt
3
4 train_data2D_1 = np.random.multivariate_normal(mean=mul1, cov=sigma1, size
                                                =size)
5 train_data2D_1_label = np.zeros((size, 1))
6
7 train_data2D_2 = np.random.multivariate_normal(mean=mul2, cov=sigma2, size
                                                =size)
8 train_data2D_2_label = np.ones((size, 1))
9 X_train = np.vstack([train_data2D_1, train_data2D_2])
10 y_train = np.vstack([train_data2D_1_label, train_data2D_2_label])
11
12 test_data2D_1 = np.random.multivariate_normal(mean=mul1, cov=sigma1, size=
                                                size)
13 test_data2D_1_label = np.zeros((size, 1))
14
15 test_data2D_2 = np.random.multivariate_normal(mean=mul2, cov=sigma2, size=
                                                size)
16 test_data2D_2_label = np.ones((size, 1))
17 X_test = np.vstack([test_data2D_1, test_data2D_2])
18 y_test = np.vstack([test_data2D_1_label, test_data2D_2_label])
19
20 pred, posterior, err = myNB(X_train, y_train, X_test, y_test)
21
22 # report the metrics
23 acc = metrics.accuracy_score(np.array(pred), y_test.squeeze())
24 print("Accuracy = %.3f" % (sum([p==a for p, a in zip(np.array(pred),
                                                y_test.squeeze())])/len(np.array(pred)
                                                ))))
25 precision = metrics.precision_score(np.array(pred), y_test.squeeze())
26 print("Precision = %.3f" % precision)
27 recall = metrics.recall_score(np.array(pred), y_test.squeeze())
28 print("Metrics = %.3f" % recall)
29 cm = metrics.confusion_matrix(np.array(pred), y_test.squeeze())
30 print("Confusion Matrix = {}".format(cm))
31
32 # draw data points
33 plt.scatter(train_data2D_1[:, 0], train_data2D_1[:, 1], label='$\mu_1, \backslash
                                                sigma_1$')

```



```

34 plt.scatter(train_data2D_2[:, 0], train_data2D_2[:, 1], label='$\mu_2, \sigma_2$')
35 plt.legend()
36 plt.show()
37 -----
38 Accuracy = 0.912
39 Precision = 0.896
40 Metrics = 0.926
41 Confusion Matrix = [[464  52]
42                    [ 36 448]]

```




Figure 8: Dataset distribution with two different colors

And in my training data, I have changed the number of examples in each class to 10,20,50,100,300,500 and perform prediction on the testing data with my code. The plot of changes of accuracies are shown below. The result shows that the more examples we have, more higher accuracy we obtain.

```

1 # change the number of examples
2 accs = []
3 x_axis = [0, 1, 2, 3, 4, 5]
4 for n in [10,20,50,100,300,500]:
5     train_set_1 = train_data2D_1[:n]
6     train_label1 = np.zeros((n, 1))
7     train_set_2 = train_data2D_2[:n]
8     train_label2 = np.ones((n, 1))
9     X_train = np.vstack([train_set_1, train_set_2])
10    y_train = np.vstack([train_label1, train_label2])
11    pred, posterior, err = myNB(X_train, y_train, X_test, y_test)
12    acc = metrics.accuracy_score(np.array(pred), y_test.squeeze())
13    accs.append(acc)

```

```

14
15
16 plt.plot(x_axis, accs, 'ro-', label='Changes of accuracies')
17 plt.xlabel("number of examples")
18 plt.ylabel("accuracy")
19 plt.legend()
20 plt.show()

```

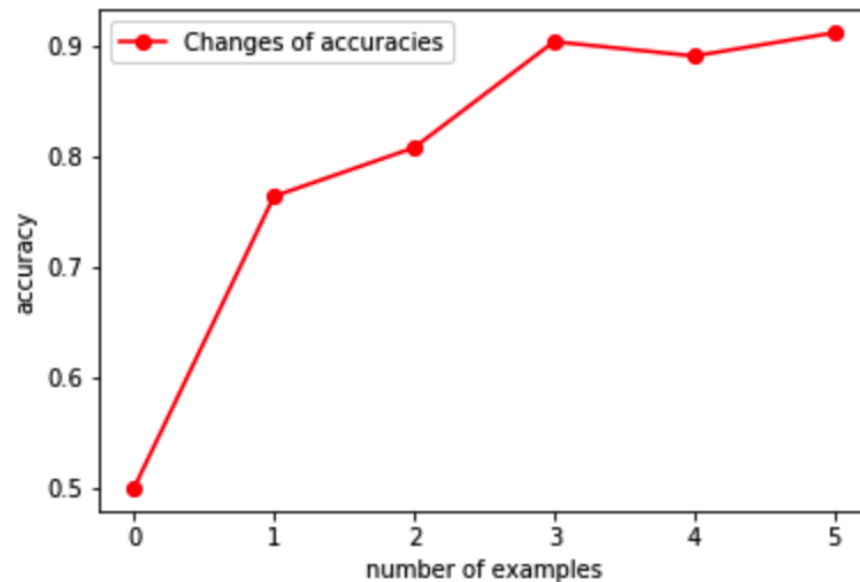


Figure 9: The changes of accuracies along with different number of examples

And when in the training data, change the number of examples in class 0 as 700 and the other as 300. Perform prediction on the testing dataset, I found that the accuracy becomes way much lower because of the unbalanced data.

```

1 # change the number of examples
2 train_data2D_1 = np.random.multivariate_normal(mean=mul1, cov=sigma1, size
                                                =700)
3 train_data2D_1_label = np.zeros((700, 1))
4
5 train_data2D_2 = np.random.multivariate_normal(mean=mul2, cov=sigma2, size
                                                =300)
6 train_data2D_2_label = np.ones((300, 1))
7
8 train_set_1 = train_data2D_1[:700]
9 train_label1 = np.zeros((700, 1))
10 train_set_2 = train_data2D_2[:300]
11 train_label2 = np.ones((300, 1))
12 X_train = np.vstack([train_set_1, train_set_2])
13 y_train = np.vstack([train_label1, train_label2])
14 pred, posterior, err = myNB(X_train, y_train, X_test, y_test)
15 acc = metrics.accuracy_score(np.array(pred), y_test.squeeze())

```

```

16 print("The accuracy when change the number of examples in class 0 as 700
    and the other as 300:", acc)
17 -----
18 The accuracy when change the number of examples in class 0 as 700 and the
    other as 300: 0.834

```

Then the ROC and AUC curve are shown below:

No Skill: ROC AUC=0.500  
 Logistic: ROC AUC=0.863




Figure 10: ROC curves from the dataset above

In the last part, I reproduce the tfidf matrix in the homework1 and print some of the reviews.

```

1 # Generate tfidf matrix
2 !pip install nltk
3 import nltk
4 nltk.download('wordnet')
5 nltk.download('stopwords')
6 import pandas as pd
7
8 review=pd.read_csv('./data/Amazon_Reviews.csv')
9 review['Label']=review['Label'].map({'__label__2 ':1,'__label__1 ':0})
10 docs_num = review.shape[0]
11
12 # Preprocessing for tokenize, stopwords removal, lemmatization
13 from nltk.tokenize import RegexpTokenizer
14 from nltk.stem import WordNetLemmatizer
15 from nltk.corpus import stopwords
16
17 tokenizer=RegexpTokenizer(r'\w+')
18 lemmatizer=WordNetLemmatizer()
19

```

```

20 for idx, row in review.iterrows():
21     tmp = []
22     tokens = tokenizer.tokenize(row['Review'])
23     for w in tokens:
24         if w.lower() not in stopwords.words('english'):
25             tmp.append(lemmatizer.lemmatize(w.lower()))
26     review.iloc[idx, 0] = ' '.join(tmp)
27
28 DF = {}
29 for idx, row in review.iterrows():
30     for w in row['Review'].split():
31         try:
32             DF[w].add(idx)
33         except:
34             DF[w] = {idx}
35 DF = {k: len(v) for k, v in DF.items()}
36 vocab = [k for k, v in DF.items()]
37 vocab_len = len(vocab)
38
39 from collections import Counter
40 import numpy as np
41
42 TFIDF = {}
43 for idx, row in review.iterrows():
44     rw = row['Review'].split()
45     cnt = Counter(rw)
46     for w in np.unique(rw):
47         tf = cnt[w]/len(rw)
48         df = DF[w]
49         idf = np.log((docs_num+1)/(df+1))
50         TFIDF[(idx, w)] = tf*idf
51
52 n = 0
53 for k, v in TFIDF.items():
54     if n<5:
55         print(k, v)
56         n += 1
57     else:
58         break
59
60 """
61 P: like(54) great(38) good(35) love(33) best(20)
62 N: however(15) disappointed(11) waste(10) poor(10) hard(8)
63 """
64 def gen_vector(idx, review):
65     Q = np.zeros(10)
66     P = np.zeros(10)
67     ten_words = ['like', 'great', 'good', 'love', 'best',
68                 'however', 'disappointed', 'waste', 'poor', 'hard']
69     for w in ten_words:
70         for k, v in TFIDF.items():
71             if k[0]==idx and k[1]==w:
72                 Q[ten_words.index(w)] = v
73     for word in review:

```

```

74         if word==w:
75             P[ten_words.index(w)] += 1
76
77     return P, Q
78
79 tfidf_vectors = []
80 count_vectors = []
81 for idx, row in review.iterrows():
82     rw = row['Review'].split()
83     P, Q = gen_vector(idx, rw)
84     tfidf_vectors.append([Q, row['Label']])
85     count_vectors.append(P)
86
87 for i in range(10):
88     print("tf-idf vector for document {} is {}".format(i, tfidf_vectors[i]
89                                                     ))
89 -----
90 tf-idf vector for document 0 is [array([0.          , 0.          , 0.
91         0.          , 0.          , 0.          , 0.          , 0.05008433,
92         0.          , 0.          , 0.          , 0.          , 0.          ]), 1]
93 tf-idf vector for document 1 is [array([0.          , 0.          , 0.
94         0.          , 0.          , 0.          , 0.          , 0.09799108,
95         0.          , 0.          , 0.          , 0.          , 0.          ]), 1]
96 tf-idf vector for document 2 is [array([0.02129965, 0.          , 0.
97         0.          , 0.          , 0.          , 0.          , 0.03266369,
98         0.          , 0.          , 0.          , 0.          , 0.          ]), 1]
99 tf-idf vector for document 3 is [array([0.05878704, 0.          , 0.
100        0.          , 0.          , 0.          , 0.          , 0.0300506 ,
101        0.          , 0.          , 0.          , 0.          , 0.          ]), 1]
102 tf-idf vector for document 4 is [array([0.          , 0.          , 0.03727823
103         0.          , 0.          , 0.          , 0.          , 0.04899554,
104         0.          , 0.          , 0.          , 0.          , 0.          ]), 1]
105 tf-idf vector for document 5 is [array([0.          , 0.          , 0.
106         0.          , 0.          , 0.          , 0.          , 0.06091338,
107         0.          , 0.          , 0.          , 0.          , 0.          ]), 1]
108 tf-idf vector for document 6 is [array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0
109         ., 0.]), 0]
110 tf-idf vector for document 7 is [array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0
111         ., 0.]), 1]
112 tf-idf vector for document 8 is [array([0.          , 0.          , 0.
113         0.          , 0.03937682, 0.          , 0.          , 0.          ,
114         0.          , 0.          , 0.          , 0.          , 0.          ]), 1]
115 tf-idf vector for document 9 is [array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0
116         ., 0.]), 1]

```

The I use tf-idf weight matrix as features and perform 5-fold cross-validation with the Naive Bayes classifier. The result (average accuracy, precision and reacall across all folds) are shown below.

```

1 from sklearn.model_selection import KFold
2 from sklearn.naive_bayes import GaussianNB
3 from sklearn import metrics
4

```

```

5 dataset = [[v[0], v[1]] for v in tfidf_vectors]
6 kf = KFold(n_splits=5)
7 avg_acc, avg_precision, avg_recall = [], [], []
8 for train_idx, test_idx in kf.split(dataset):
9     gnb = GaussianNB()
10    X_train, y_train = [dataset[i][0] for i in train_idx], [dataset[i][1]
11                                                              for i in train_idx]
12    X_test, y_test = [dataset[i][0] for i in test_idx], [dataset[i][1] for
13                                                         i in test_idx]
14    pred = gnb.fit(X_train, y_train).predict(X_test)
15    acc = metrics.accuracy_score(np.array(pred), y_test)
16    print("Accuracy = %.3f" % (sum([p==a for p, a in zip(np.array(pred),
17                                                         y_test)])/len(np.array(pred))))
18    precision = metrics.precision_score(np.array(pred), y_test)
19    print("Precision = %.3f" % precision)
20    recall = metrics.recall_score(np.array(pred), y_test)
21    print("Recall = %.3f" % recall)
22    print('-----'*4)
23    avg_acc.append(acc)
24    avg_precision.append(precision)
25    avg_recall.append(recall)
26
27 print("Average Accuracy is: ", sum(avg_acc)/5)
28 print("Precision Accuracy is: ", sum(avg_precision)/5)
29 print("Recall Accuracy is: ", sum(avg_recall)/5)
30
31 -----
32
33 Accuracy = 0.675
34 Precision = 0.955
35 Recall = 0.636
36 -----
37 Accuracy = 0.625
38 Precision = 1.000
39 Recall = 0.583
40 -----
41 Accuracy = 0.700
42 Precision = 0.913
43 Recall = 0.677
44 -----
45 Accuracy = 0.725
46 Precision = 1.000
47 Recall = 0.667
48 -----
49 Accuracy = 0.641
50 Precision = 1.000
51 Recall = 0.622
52 -----
53 Average Accuracy is: 0.6732051282051282
54 Precision Accuracy is: 0.9735177865612649
55 Recall Accuracy is: 0.6370809225647934

```

The result shows that cross validation can reduce bias and every data points get to be tested exactly once and is used in training  $k-1$  times. The variance of the resulting estimate is reduced as  $k$  increases.

## Reference

- <https://jakevdp.github.io/PythonDataScienceHandbook/05.13-kernel-density-estimation.html>
- <https://stats.stackexchange.com/questions/206963/adaptive-variable-bandwidth-in-2d-gaussian-kernel-density-estimator-to-account>